

# Transforming malicious code to ROP gadgets for antivirus evasion

ISSN 1751-8709  
 Received on 26th July 2018  
 Revised 22nd February 2019  
 Accepted on 13th March 2019  
 E-First on 3rd June 2019  
 doi: 10.1049/iet-ifis.2018.5386  
[www.ietdl.org](http://www.ietdl.org)

Christoforos Ntantogian<sup>1</sup>, Georgios Poulios<sup>1</sup>, Georgios Karopoulos<sup>2</sup>, Christos Xenakis<sup>1</sup> 

<sup>1</sup>Department of Digital Systems, University of Piraeus, Piraeus, Greece

<sup>2</sup>Department of Informatics and Telecommunications, University of Athens, Athens, Greece

E-mail: xenakis@unipi.gr

**Abstract:** This study advances research in offensive technology by proposing return oriented programming (ROP) as a means to achieve code obfuscation. The key inspiration is that ROP's unique structure poses various challenges to malware analysis compared to traditional shellcode inspection and detection. The proposed ROP-based attack vector provides two unique features: (i) the ability to automatically analyse and generate equivalent ROP chains for a given code, and (ii) the ability to reuse legitimate code found in an executable in the form of ROP gadgets. To this end, a software tool named ROPInjector was developed which, given any piece of shellcode and any legitimate executable file, it transforms the shellcode to its ROP equivalent re-using the available code in the executable and finally patches the ROP chain infecting the executable. After trying various combinations of evasion techniques, the results show that ROPInjector can evade nearly and completely all antivirus software employed in the online VirusTotal service, making ROP an effective ingredient for code obfuscation. This attack vector poses a serious threat which malicious actors can take advantage to perform cyber-attack campaigns.

## 1 Introduction

Return oriented programming (ROP) gained increased attention during the late 2000s [1] as an advanced stack smashing method to bypass security mechanisms. ROP is a rediscovery of threaded code in which programs typically consist of a chain of addresses in the stack pointing to code chunks in the attacked executable or its loaded libraries each of them ending with a return instruction. These borrowed code chunks are called *gadgets* and their 'return' is in fact a call to the next gadget in the chain. As an analogy to regular code, in ROP, gadgets are the 'instructions' and esp is the program counter.

Traditionally, the primal usage scenario of ROP has been (and still is) software exploitation to bypass non-executable stack and heap. Other recent works have proposed ROP for other purposes such as rootkit development [2], software watermarking [3], steganography [4], and code integrity verification [5]. In this paper, we argue that ROP has another field of application which is code obfuscation and antivirus (AV) evasion. We leverage the fact that ROP is Turing complete and therefore it can perform arbitrary computations; thus every shellcode can be transformed to a ROP equivalent. Our intuition is that ROP's unique structure poses various challenges compared to traditional shellcode detection including verbosity of the gadgets, stack-based chaining, lack of immediates, and the distinction of function calls and regular control flow. On top of that, ROP inherently provides polymorphic properties by selecting a different equivalent gadget for transforming a given instruction or by chaining the selected gadgets in a different order. Another benefit of using ROP for AV evasion is that such borrowed code (that of gadgets) is always benign and tested against false positives. Thus, considering the aforementioned challenges, at its core, ROP can be seen as a new method for code obfuscation.

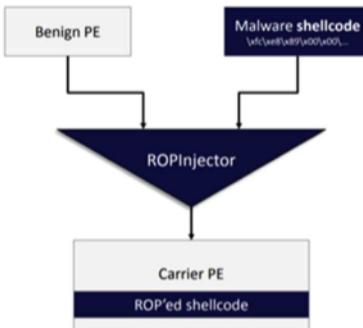
Based on the above observations, in this paper, we present ROPInjector, a tool which, given any piece of code (hereafter, also referred to as shellcode) and any non-packed executable file, it transforms the shellcode into its ROP equivalent and patches it into (i.e. infects) the executable file. ROPInjector [6], which is implemented in C/C++ programming language, infects Portable Executables (PEs) for Windows OS. Since it is very common for AVs to detect minor deviations from the typical arrangement of the

file sections and their characteristics (e.g. a second executable section with RWX permissions), besides the transformation of the code into a non-recognisable, non-recurrent form, the developed tool addresses several additional issues to achieve evasion, such as the positioning of the shellcode in the carrier executable and the way of transferring control to the shellcode. Moreover, we have performed several experiments to evaluate the effectiveness of the proposed tool by injecting shellcodes to well-known executable files including acrobat reader, firefox, and so on. The high evasion results show that ROPInjector combined with simple behavioural anti-profiling techniques can render AV ineffective and incapable to detect ROP-based malware.

The rest of the paper is organised as follows. Section 2 presents and compares with ROPInjector the related work. Section 3 elaborates on the architecture of the proposed ROPInjector and its functionality details. In Section 4, we analyse experimental results and in Section 5 we provide a discussion of possible mitigation techniques and the challenges of ROPInjector. Finally, Section 6 concludes the paper.

## 2 Related work

While the traditional use of ROP is software exploitation (i.e. bypass non-executable stack and heap), there are some previous works that have proposed alternative uses of ROP. More specifically, in [3] the authors propose ROP for benign purposes; specifically, they use ROP for software watermarking. The proposed ROP-based watermarking is able to transform watermarking code into ROP gadgets and build them in the data region. Once triggered using a secret message, the pre-constructed ROP execution will recover the hidden watermark message. The proposed method ensures that the watermarked program does not have an explicit code stream that belongs exclusively to watermarking. Instead, the authors use operating system libraries to borrow the ROP gadgets, preventing detection by software analysis. Towards this direction, RopSteg [4] has been proposed for program steganography. The latter is a variation of software obfuscation but it differs from it, since in program steganography the instructions are hidden instead of being transformed. RopSteg achieves to hide selected code protection by generating equivalent ROP gadgets and blending them into the executable. Both works



**Fig. 1** ROPInjector high-level functionality

[3, 4] are reminiscent of ROPInjector as they take advantage of ROP to hide a message or a piece of code inside a PE. However, none of these two solutions evaluate the detection capabilities of ROP, as their main intention is the use of ROP for benign purposes such as software watermarking and steganography. In addition, these solutions do not release a tool to further evaluate their capabilities. Finally, ROPOB [7] is a code transformation technique to obfuscate control flow using ROP. The main contribution of ROPOB is that due to the use of ROP to complete control flow transfer, static reverse engineering methods cannot easily discover the real control flow of the program. However, the main limitation of ROPOB is that through dynamic analysis the obfuscation trivially breaks.

A similar work to ROPInjector is presented in [8], in which the authors leverage ROP to hide malicious code in a PE file. The novel idea is to encrypt the ROP chain to avoid detection and remotely trigger the ROP execution at the runtime. Although this idea is interesting, the encryption itself can raise suspicions, due to high entropy, while a proper key management solution always requires extra effort. Nonetheless, the solution described in [8] presents yet another interesting use case of ROP for malicious purposes. The work closest to ours is presented by Mohan and Hamlen [9]. The authors have developed a metamorphic obfuscator called Frankenstein which is able to reassemble a given malware with code fragments entirely from other benign programs. Authors' motivation was the creation of malware variations from benign pieces of entirely randomly selected binaries residing in a system to avoid signature matching detection. They deduce the problem of generating mutations into a searching problem. Their proposed method is able to search for segments of code found in benign binaries and evaluate them semantically with the given malicious instructions. Finally, it performs the suitable code arrangements to construct the final payload. Frankenstein differs from our proposed ROPInjector in several aspects. First, Frankenstein considers a relaxed version of a gadget. That is, while in the classic definition of ROP [10], a gadget ends with the *ret* instruction, Frankenstein definition of a gadget is any sequence of bytes that are interpretable as valid x86 instructions. On the contrary, ROPInjector relies on the pure and original definition of ROP gadgets. Second, Frankenstein main purpose is to modify only code snippets which may look suspicious, while ROPInjector transforms the whole shellcode to the ROP counterpart. Third, authors have only implemented a proof of concept (which is not available publicly).

Finally, we examine two non-academic works that have the same purpose with ROPInjector; that is, to infect PE files with common (possibly encrypted) shellcode in order to bypass AV software.

The first tool named Shellter [11] focuses on maintaining the original structure of the PE file, by avoiding injection of the shellcode into predefined locations or changing the characteristics of the existing sections. It achieves so by overwriting existing code for which it is certain that will be given control during execution of the program. The latter is deduced by tracing the executable file and analysing its execution flow. Shellter is also capable of reusing imports of the original PE file to change the writing permissions of the section containing the shellcode so that encrypted and self-modifying code can be used. It is also capable of injecting 'junk code' before the shellcode that delays execution as a means to anti-

emulation. Shellter is advanced in terms of dynamically selecting the location of the patch in the shellcode (as opposed to extending the text section). However, while it features a patching method that introduces variability (as to where in the file is the shellcode injected), it relies on traditional polymorphism methods, that are still subject to signature generation and detection of write permissions or modifications of the text section in memory. Moreover, our proposed approach can introduce variability too, due to the transformation to ROP (which is dependent on the PE file).

The second tool is PEinject [12], which is mostly a method (and referenced as such) rather than a full-featured tool. It injects the shellcode in the (first sufficiently large) padding space found in the text section (either 0xCC caves or section padding) and does not encode or modify the payload in any way, neither does it anticipate for self-modifying or encrypted payloads. Control is passed to the injected shellcode by modifying the address of entry point of the PE file's NT\_HEADER.

A common limitation of the aforementioned related works is the limited applicability and practicability of their proposed techniques. For instance, Borrello *et al.* [8] discuss only theoretically the proposed solution, while Lu *et al.* [4] present a proof of concept which has been manually prepared only for evaluation purposes. On the other hand, Mohan and Hamlen [9] apply the proposed techniques only to a very limited number of code snippets. In contrast to all other previous works, ROPInjector is the first tool released as open source that exploits the polymorphic properties of ROP for advanced code obfuscation. ROPInjector has been designed to be generic and scalable in the sense that given as an input any functional shellcode and any legitimate executable, it will produce an executable that will retain its legitimate functionality bundled with the malicious shellcode. Finally, ROPInjector is fully automated and no manual work is required.

### 3 ROPInjector

#### 3.1 High-level overview

ROPInjector takes as an input a benign PE and a malicious piece of code (i.e. shellcode) and outputs the PE infected with the malicious piece of code in a ROP form (see Fig. 1). In general, the ROPInjector approach can be divided into five distinct phases as follows:

- i. Reverse analysis of machine code
- ii. Finding ROP gadgets in PE
- iii. ROP compilation and injecting gadgets
- iv. Chaining gadgets
- v. Passing control to the shellcode

These phases are executed sequentially one after the other. If any of the above phases fail to complete, then ROPInjector outputs an error message describing the cause of the failure. In the next sections, we are going to analyse each of the above phases providing detailed examples to gain a better understanding of the presented notions.

#### 3.2 Reverse analysis of machine code

Reverse analysis of machine code into data structures that are easy to handle is crucial to perform any kind of patching, modifications, re-assemble, and any transformation to ROP. Two are the most important pieces of information required: (i) the origin and destination of all relative references (e.g. a relative jump and its target) and (ii) which registers are being written or read during each instruction, as well as which registers are free to modify. The former is required for injecting or removing instructions from a code segment without breaking its functionality. The latter is particularly useful to enhance gadget matching, either by performing permutations, or by using gadgets that contain redundant but safe instructions (in this case, unsafe are branch, privileged, or indirect addressing mode instructions because they risk raising errors such as access violation).

### 3.2.1 MOD/REG/RM and scaled index byte (SIB)

**unrolling:** Instructions using the MOD/REG/RM indirect addressing mode with displacement or the SIB addressing scheme in the shellcode are treated specially before the transformation to ROP. Such instructions are unwanted for the following reasons:

- i. They are long (in the best and not so likely case 3 bytes long: 1 for opcode, 1 for MOD/REG/RM and 1 for SIB) hence unlikely to be found in gadgets;
- ii. They often read many general purpose registers at once, thus reserving them while as mentioned earlier, the more the free registers the better;
- iii. Their respective gadgets (should they be found or injected) will probably not be reusable, due to the use of displacement and index constants (e.g. `mov edx, [esi*2 + 16]`).

In order to circumvent this kind of situations, we reduce such instructions to their arithmetic equivalents one-by-one. We call this process *unrolling* and it is performed to the shellcode before any transformation to ROP. For instance, `mov eax, [ebx + ecx*2]` may be replaced by

```
[a'] Mov eax, ecx
[b'] sal eax, 1
[c'] add eax, ebx
[d'] mov eax, [eax]
```

If the register *eax* is not free to use for the arithmetic operations, another temporary register that is free may be used.

Noteworthy is how unrolling unlocks register access from one atomic instruction to many. For instance, in the example, *ecx* is freed at [a'] and *ebx* at [c']. If, for example, *eax* were to be freed at the preceding 10 instructions, then instructions [a'] to [c'] could be moved 10 instructions behind, thus resulting in an additional free register (i.e. *ecx* and *ebx*, but not *eax* which will not be free) in that preceding code chunk.

### 3.3 Finding gadgets

Candidate gadgets in the executable sections of the given PE file must end in one of the following instructions:

- `ret`,
- `retn`,
- `pop regX`,
- `jmp regX`, or `jmp regX`.

Exceptionally for the latter, the gadget in question must be first paired with a *loader gadget* that loads the required return address into *regX*. The process begins by finding all gadget endings and temporarily storing them to a list. For each of those endings, *n* bytes of preceding machine code are disassembled for each *n* up to maximum depth *N* (typically 20 bytes). If such disassembly aligns with the ending (not guaranteed since x86 instructions are of variable length) a candidate gadget has been found. Candidate gadgets containing any illegal, privileged (e.g. *sysenter*, *int*, *iret*), branch or *esp* modifying instructions are filtered out.

### 3.4 ROP compilation and injecting gadgets

The gadgets found in the aforementioned process are first analysed instruction-by-instruction to infer register access. Since gadgets are allowed to contain safe but redundant instructions, their register access is tested for modifications to the register in question (e.g. `mov ecx, eax; pop ecx; ret`; gadget cannot be used for moving *eax* to *ecx*) as well as the non-free registers of the source instruction to be encoded.

Following that, they are lifted into an intermediate representation (IR) consisting of an operation-type, and three operands with different meaning depending on the type. If a multi-instruction gadget contains more than one representable instructions, only the first is considered. However, the following ones have also been considered in other gadgets with the same ending, because of the backwards gadget finding process described in the previous paragraph. Noteworthy is the fact that by parsing into this higher level IR, one-to-one permutations are automatically performed. That is because both gadgets and instructions are classified into one of these types, based on which the encoding is then performed, rather than on the instructions per se. The IR is also useful for selecting the encoder function accompanying every gadget. Encoders are responsible to answer ‘whether their assigned gadget can encode a given instruction’, as well as to encode it into a list of stack operations if requested to.

Predefined, one-to-one permutations (i.e. one instruction to one gadget) are achieved through the IR and encoder functions. Encoders will also perform basic algebraic permutations based on the properties of addition, subtraction, multiplication, and division. For instance, if the instruction to be encoded is of type ADD\_IMM (*add reg, imm*), an encoder will repeat anything *add reg, x* with *x* being an integer divisor of *imm*, *imm/x* times. Addition and subtraction with constants will also be swapped if the signs of the constants are flipped. *M*-to-*N* permutations quickly scale to exponentially growing space and are out of the scope of this work.

In order to enhance transformation of the source shellcode, and since not all required gadgets are always found in the PE file, new ones are also injected as needed. Firstly, the 0xCC caves (i.e. padding commonly left by the linker in-between code segments in the text section of PEs) are used for this injection, and if they are filled, the text section is extended before the actual patch. The injection is performed in the least noticeable way to avoid alarms. If a standard epilogue (`mov esp, ebp; pop ebp; ret`) is found right before the 0xCC cave, the gadget is injected in-between the preceding code and the epilogue. Fig. 2 depicts such an example gadget injection of a `mov ecx, eax` gadget. In the case that no epilogue is found at the boundary with the 0xCC cave, a pseudo-function with standard prologue and epilogue is injected to avoid heuristics or *n*-grams that might raise suspicion due to non-ordinary returns. This pseudo-function has the following form shown in Fig. 3.

Algorithm 1 as shown in Fig. 4 is used when ROPInjector builds the ROP equivalent of a given shellcode. The first step of the algorithm is to scan the PE and find all the pieces of code that could potentially be used for chaining to the provided shellcode, named gadget endings. As a last step of the initialisation, the algorithm locates the 0xCC nests (also known as caves) where gadgets can be generated and injected if not found in the PE. In case one or more suitable gadgets are found for the instruction, then the best one is chosen. In this scope, the best gadget is

|   |  |
|---|--|
| <pre>mov esp, ebp pop ebp ret(n) CCCCCCCCCCCCCCCCCCCC</pre> | <pre>jmp epilogue; normal flow avoiding gadget mov ecx, eax; the injected gadget jmp return; gadget flow avoiding std. epilogue  epilogue:     mov esp, ebp     pop ebp  return:     ret(n)     CCCCCCCC</pre> |
|---|--|

Fig. 2 Injection of gadget (right) in 0xCC cave preceded by standard function epilogue (left)

```

push ebp
mov ebp, esp
<gadget code>
jmp return
mov esp, ebp
pop ebp
return:
ret

```

**Fig. 3** Pseudo-function ending used during gadget injection

**Algorithm 1** ROP compilation

**Input:** shellcode, PE

**Output:** gadgets\_list

```

1: count gadget endings
2: find CCnests in PE
3: i ← first instruction of shellcode
4: while not end of instructions do
5:   if i can't be ROP-compiled then
6:     continue
7:   else
8:     gadget ← best matching gadget
9:     if best gadget not found in PE then
10:      gadget ← inject gadget in CC nests
           // Described in Algorithm 2
11:    end if
12:  end if
13:  if gadget not targeted by jump then
14:    chain to the previous gadget
15:  else
16:    set as new gadget chain
17:  end if
18:  append gadget to gadgets_list
19:  i ← next instruction
20: end while
21: return gadgets_list

```

**Fig. 4** ROP compilation algorithm

**Algorithm 2** Injecting gadgets in CC nests

**Input:** instruction

**Output:** gadget

```

1: gadget ← assemble new gadget for instruction
2: cc_nest ← find space for gadget
3: if not enough space then
4:   cc_nest ← make space from padding or extending .TEXT
5: end if
6: if cc_nest is a new gadget section then
7:   write function prologue
8:   inject gadget
9:   write function epilogue
10: else
11:   add gadget to existing section
12: end if
13: return gadget

```

**Fig. 5** Gadget injection algorithm

considered to be the one that contains less instructions. Next, all instructions are checked against the list of candidate gadgets. We have to distinguish one important case here: in case an instruction cannot be encoded by any of the gadgets found in the previous step, then a suitable gadget needs to be generated and injected (see line 10 of Algorithm 1). For this purpose, the 0xCC nests are considered as presented in Algorithm 2 (see Fig. 5). That is, if there is space, then the needed gadget is injected in the caves. On the other hand, if the gadget cannot fit, then the text section of the PE is extended to create more space.

Continuing now with Algorithm 1, gadgets either injected or found, are used to encode each one of the (allowed) instructions contained in the shellcode. The last step of this algorithm concerns the chaining of the gadgets (either found or injected). For this part we have to distinguish two cases:

- If a gadget is targeted by a branch instruction, then the algorithm handles it as a new chain where more gadget chains can be added.
- If the gadget is not targeted, the algorithm just adds it to an existing gadget chain.

### 3.5 Chaining gadgets

The return address chain can be built either during runtime or during compile-time and saved to the initialised data section of the file (to be then copied at runtime to the stack). The most alarming option would be the first (during runtime) and we choose this to evaluate our evasion ratio. During this process, besides the pushing of the VAs onto the stack, the ROP compiler must consider pushing immediate constants, adjustments for stack pointer modifications in the gadget (e.g. redundant *rops*, *retns*) and gadgets with loader gadgets. For this purpose, the following types of *stack operations* are defined:

PUSH\_VA; push a (loader) gadget VA onto the stack

PUSH\_IMM; push an immediate constant onto the stack

ADVANCE; advance (subtract from) the stack pointer a number of bytes

CHAIN; pseudo-operation denoting a placeholder for the next gadget's VA

The result of the encoding process of a given instruction by a given gadget is a series of stack operations for the invocation of the gadget. The list of such operations for all gadget calls describes the assembly instructions that if executed, will build the chain in the stack. Alternatively, such operations may be used to create the required stack frame during compile-time, save it as initialised data and copy it over from the data section during runtime. The latter process allows also for encoding/decoding of the stack frame. In the former case, and when multiple calls are made to the same gadget (e.g. as in using *inc eax* to achieve *add eax, 3*) the compiler wraps the call with a conditional jump loop using a free register.

However, not all types of instructions can be easily encoded into ROP. In this work, we do not consider the encoding of branches (jumps, calls, loops, interrupts), privileged instructions, and pops. Hence, the return-oriented code chunks must finally return back to the source shellcode. This is achieved by wrapping the chain building instructions in the following:

```

[1] call build_chain
[2] jmp past_the_chain
[3] build_chain:
[4] push <VA of gadget N>
[5] ....
[6] push <VA of gadget 1>
[7] ret
[8] past_the_chain:
[9] <other instructions/chains>

```

In this way, the last gadget (*N*) will return to instruction [12] jumping past the chain building instructions and continuing normal execution flow.

### 3.6 PE patching and passing control to the shellcode

First of all, the patching of the PE file and the passing of control to the shellcode must be done in the least noticeable way. A second executable section hosting the shellcode would be too alarming, since the vast majority of executables has only one. The next least disruptive and easy to implement option would be to inject the shellcode once again in the 0xCC nests. However, there may not always be sufficient space in those 0xCC caves, as ROPIInjector puts this padding space to inject the required missing gadgets.

For the above reasons, we choose to append the shellcode to the existing text section of the executable, and correct all section headers and relocations accordingly. To pass control to it, the default practice is to replace the instructions pointed to by

NT\_HEADER.AddressOfEntryPoint with a jump to the shellcode, which is appended to those replaced instructions followed by a jump back to the original execution flow. Directly pointing the address of entry point to the shellcode in this case is avoided, since many AVs' heuristics are alarmed by the fact that it points towards the end of text section. An alternative to giving control to the shellcode at program entry is to hook any calls to ExitProcess, exit or other similar functions. This technique in particular, as shown also later by the results, bypasses behavioural profiling by AVs that employ emulation or sandboxing. This can be attributed to the fact that either AVs emulate only a small portion of the executable's entry code due to scanning time constraints, or because of a lack of techniques for triggering a graceful exit (i.e. many programs do not handle SIGINT and SIGTERM signals).

## 4 Experiments and results

### 4.1 Methodology

In this section, we present the methodology adopted to evaluate ROPInjector. We have selected for carrier PEs 9 popular 32-bit executables (see first column of Table 1). Regarding the source shellcode, we selected two of the most popular payloads of Metasploit framework [13]: (i) (staged) meterpreter reverse tcp, and (ii) the reverse tcp shell. For each PE (nine in total) and each shellcode (two in total) we performed four patching scenarios as analysed in Table 2: (i) ROP entry, (ii) ROP exit, (iii) Shellcode entry, (iv) Shellcode exit. Thus, we created 72 different samples (i.e. 72 infected PEs). On these samples we derived the following results:

- Various statistics including: (i) number of candidate gadgets found in the PEs, (ii) number of gadgets re-used from the PEs, (iii) missing gadgets that were injected, (iv) number of instructions of the two tested shellcodes that were transformed to ROP equivalents.
- Similarity measurements between original PE and infected PE.
- ROPInjector Code complexity metrics.
- Evasion results for each PE and top 10 popular AV detection results.
- Evasion results for Enhanced Mitigation Experience Toolkit (EMET).

Note that all conducted experiments were performed on a Windows 7 32-bit operating system.

**Table 1** List of PE files used as carriers in the experiments and various statistics (ROP entry method and reverse tcp shell)

| PE            | Size, kB | File version      | Candidate gadgets found in PE | Gadgets injected | Gadgets used from PE |
|---------------|----------|-------------------|-------------------------------|------------------|----------------------|
| Acrobat.exe   | 650      | 19.10.20069.49826 | 338                           | 78               | 12                   |
| AcroRd32.exe  | 1423     | 11.0.8.4          | 5599                          | 67               | 26                   |
| cmd.exe       | 305      | 6.3.9600.16384    | 608                           | 76               | 16                   |
| firefox.exe   | 439      | 63.0.3.6892       | 1634                          | 65               | 28                   |
| java.exe      | 187      | 8.0.192.12        | 1148                          | 76               | 18                   |
| nam.exe       | 1828     | 1.0a11a           | 3515                          | 67               | 25                   |
| notepad++.exe | 2783     | 7.6.0.0           | 7778                          | 55               | 41                   |
| Rainmeter.exe | 39       | 2.4.0.1678        | 11                            | 83               | 0                    |
| wmplayer.exe  | 163      | 12.0.9600.19145   | 60                            | 82               | 3                    |

**Table 2** List of patching methods tested

| Patch method                       | Description  |
|------------------------------------|--|
| ROP entry                          | the executable file is patched with the shellcode unrolled, converted to ROP and the entry point before the original PE code                                   |
| ROP exit                           | the executable file is patched with the shellcode unrolled, converted to ROP and the entry point before the original program's exit (hook ExitProcess or exit) |
| Shellcode entry (no-rop no-unroll) | the executable file is patched with the shellcode intact and the entry point before the original PE code   |
| Shellcode exit (no-rop no-unroll)  | the executable file is patched with the shellcode intact and the entry point before the original program's exit (hook ExitProcess or exit)                     |



interesting observation is that many of the popular commercial solutions used widely in organisations fail to score high in the detection of the infected PEs. For instance, McAfee failed to detect any of the infected executables, while AVG detected only one infected PE. Kaspersky and ZoneAlarm had the best performance as they detected five out of nine PEs.

#### 4.5 EMET evasion

Apart from AVs, we evaluated the evasion effectiveness of ROPInjector against a sophisticated anti-malware tool named EMET. The latter is developed by Microsoft aiming to mitigate exploitation attacks including ROP which are leveraged by malware. During the runtime of an executable, 50 security-wise critical functions are monitored by EMET (i.e. hooked) and several security checks are performed before a monitored critical function is allowed to be executed. The most important and relevant to ROP security checks of EMET are briefly discussed below:

- *Caller Check* rule ensures that when a critical function is invoked, it is reached via a Call instruction rather than a Ret.
- *SimExecFlow* is the reverse of Caller Check rule. After a call to a critical function is completed, this protection simulates the execution to ensure that the following code is legitimate (and not a ROP chain). The first return address is given on the stack, and subsequent return addresses are deduced by simulating instructions that modify the stack/frame pointer. The SimExecFlow rule checks that each return address is preceded by a Call instruction to appear legitimate.
- *StackPivot*. Upon entering a critical function, EMET checks to ensure that the stack pointer register is within the thread's upper and lower specified stack limit. This guards against pivoting the stack pointer to attacker-controlled memory such as heap.
- *MemProt* check is triggered when a specific critical function named VirtualProtect() is called, in order to check whether the attacker is changing execution permissions on stack memory pages.
- *LoadLibrary* rule ensures that a shellcode does not use remote file paths by invoking the critical function LoadLibrary().

If any of these security checks detects a ROP-based attack, EMET closes the application and creates an appropriate log event.

To evaluate the evasion effectiveness of ROPInjector against EMET detection capabilities, we scanned the ROPInjector infected executables against EMET (version 5.5). As we observed in our experiments, EMET techniques proved to be ineffective against ROPInjector as it was not able to detect any executable generated by ROPInjector. This can be attributed to the fact that the exploitation scenarios that EMET tries to protect are different from the exploitation scenarios of ROPInjector. More specifically, the aforementioned EMET rules aim at protecting against malware that try to exploit a vulnerability and perform a memory corruption attacks (i.e. buffer overflows). In such scenarios, a malware leverage ROP attacks after a buffer overflow, in order to call specific functions which would further extend its capabilities. On the other hand, ROPInjector does not rely on memory corruption attacks for its nefarious purposes; instead ROPInjector stores (i.e. backdoors) a malicious piece of code (i.e. shellcode) in a ROP-based form inside a legitimate executable. When the legitimate executable is executed (using, e.g. social engineering or HID attacks), the backdoored shellcode is also executed.

To exemplify, we consider the MemProt rule and how it fails to protect against ROPInjector. A common goal of malware is to enable a memory region to be executable, in order to execute subsequently a malicious shellcode. This is typically achieved by performing a buffer overflow to pass the control to a ROP chain, in order to call the VirtualProtect() function. EMET's MemProt rule can protect against such attacks as it monitors the VirtualProtect() function, but not against ROPInjector, since the latter neither has to call VirutalProtect() nor it has to perform memory corruption attacks as the whole executable is under its control. As a side note, it is important to mention that various research works such as the one presented in [17] have shown that EMET rules can be

bypassed trivially and in general they are insufficient to protect against memory corruption attacks.

## 5 Discussion

*Detection:* Most AV software relies on string signatures and mild behavioural profiling detection mechanisms. By encoding the shellcode into its return-oriented equivalent and even by performing elementary mutations (unrolling), the former can be bypassed in the vast majority of cases. Of course, the return address chain has to be built somehow in the stack and that would leave a footprint subject to signing. The process involves either pushing the return addresses to the stack or just copying the whole chain from another memory location (possibly some data segment) and adjusting the stack pointer. However, (a) the kind of code required for such operations is very common and seemingly benign, (b) it largely depends on the attacked PE and its image base since in the worst case it is a series of `push <VAi>` operations, and (c) can be randomised or encrypted in many and trivial ways. Especially, what we mentioned for reason (b) holds because gadget addresses change for different PEs and different image bases, hence changing the footprint and statistics of the chain building instructions even if they are for the same source shellcode. Dynamic analysis and behavioural profiling can also be avoided by carefully intercepting normal execution flow in points that AVs either cannot emulate or simply cannot derive enough evidence to classify the behaviour as malicious. In this work, we presented as a means to the latter the hooking of common calls to process exit, resulting in many cases in absolute evasion and in others rates >98%. If the anti-dynamic analysis protection of ROPInjector is bypassed, then the dynamic analysis of an infected PE can reveal the malicious code hidden in a ROP form. However, we consider this as out of scope since ROPInjector is mainly designed as an obfuscation tool to avoid static detection. Advanced techniques specialised for anti-dynamic analysis such as malwash [18] can be combined and further enhance evasion capabilities of ROPInjector.

The techniques presented can still be mitigated if dealt with individually. For instance, signatures could be created for ROP building instructions and behavioural analysis could be also performed backwards in terms of process life-cycle. However, since slight variations and randomisation can again disarm scanners, a more robust countermeasure does not seem straightforward to design, and/or practical to implement.

Several AVs rely on code statistics (such as entropy, n-grams and more), in order to classify PE files as benign or malicious. Unless such methods are designed to consider separate parts of the file's code, ROPInjector does not affect the statistic metrics of the binary file as a whole. The only metric that is affected and can be used by AVs for detection is the hash of the PE itself. In particular, for well-known PEs (such as Firefox.exe), AVs can cross-check the hash of the PE. If the hash is different from the official version of the executable, then AVs can raise an alarm. However, since new versions of these applications are continuously released, maintaining the legitimate hash values of each PE is a challenging task.

*Control flow integrity (CFI) and whitelisting:* A mechanism that counteracts against ROP attacks is CFI [19]. The latter is a compile time mechanism for preventing malicious code from redirecting the execution flow of a program. In other words, the goal of CFI is to restrict the set of possible control-flow transfers to those that are strictly required for correct program execution. This prevents code-reuse techniques such as ROP from working because they would cause the program to execute control-flow transfers which are illegal under CFI [20]. By its own definition, we can deduce that ROPInjector does not violate CFI policies, since the purpose of CFI is to defend against threats which will modify the *runtime operation* of an executable. On the contrary, ROPInjector changes the execution flow of an executable in the *compile time* and specifically, changes the entry or exit point of the executable in order to transfer the code execution into the ROP chain.

Perhaps the most promising direction is towards the strict coupling of the host operating system with the trusted software certificates (or checksums) and a 'default distrust all' policy, i.e.



- [7] Mu, D., Guo, J., Ding, W., *et al.*: 'ROPOB: obfuscating binary code via return oriented programming'. Int. Conf. on Security and Privacy in Communication Systems, Cham, 2017
- [8] Borrello, P., Coppa, E., D'Elia, D.C., *et al.*: 'The ROP needle: hiding trigger-based injection vectors via code reuse'. 34th ACM/SIGAPP Symp. on Applied Computing, Limassol, Cyprus, April 2019
- [9] Mohan, V., Hamlen, K.W.: 'Frankenstein: stitching malware from benign binaries'. USENIX Workshop on Offensive Technologies (WOOT 2012), Bellevue, WA, USA, 2012, pp. 77–84
- [10] Roemer, R., Buchanan, E., Shacham, H., *et al.*: 'Return-oriented programming: systems, languages, and applications', *ACM Trans. Inf. Syst. Secur.*, 2012, **15**, (1), p. 2
- [11] Shellter project. Available at <https://www.shellterproject.com>, accessed 15 April 2018
- [12] Injecting Shellcode into a Portable Executable (PE) using Python. Available at <http://www.debasish.in/2013/06/injecting-shellcode-into-portable.html>, accessed 15 April 2018
- [13] Metasploit. Available at <http://www.metasploit.com/>, accessed 15 April 2018
- [14] Karnik, A., Goswami, S., Guha, R.: 'Detecting obfuscated viruses using cosine similarity analysis'. First Asia Int. Conf. on Modelling & Simulation (AMS' 07), Phuket, Thailand, 2007
- [15] VirusTotal. Available at <https://www.virustotal.com>, accessed 15 June 2018
- [16] The Best Antivirus Protection for 2019. Available at <https://www.pcmag.com/roundup/256703/the-best-antivirus-protection>
- [17] DeMott, J.: 'Bypassing EMET 4.1', *IEEE Secur. Priv.*, 2015, **13**, (4), pp. 66–72
- [18] Ispoglou, K., Payer, M.: 'malWASH: washing malware to evade dynamic analysis'. WOOT, Austin, TX, USA, 2016
- [19] Abadi, M., Budiu, M., Erlingsson, U., *et al.*: 'Control-flow integrity'. Proc. of the 12th ACM Conf. on Computer and Communications Security, Alexandria, VA, USA, 2005
- [20] Burow, N., Carr, S., Nash, J., *et al.*: 'Control-flow integrity: precision, security, and performance', *ACM Comput. Surv.*, 2017, **50**, (1), p. 16
- [21] Applocker. Available at <https://docs.microsoft.com/en-us/windows/security/threat-protection/windows-defender-application-control/applocker/applocker-overview>, accessed 15 February 2018
- [22] Das, S., Werner, J., Antonakakis, M., *et al.*: 'Sok: the challenges, pitfalls, and perils of using hardware performance counters for security'. 2019 IEEE Symp. on Security & Privacy (SP), San Fransisco, CA, US, 2019, pp. 345–363
- [23] Das, S., Chen, B., Chandramohan, M., *et al.*: 'ROPSentry: runtime defense against ROP attacks using hardware performance counters', *Comput. Secur.*, 2018, **73**, pp. 374–388
- [24] Wang, X., Backer, J.: 'SIGDROP: signature-based ROP detection using hardware performance counters', arXiv preprint arXiv:1609.02667, 2016
- [25] Tang, A., Sethumadhavan, S., Stolfo, S.J.: 'Unsupervised anomaly-based malware detection using hardware features'. Int. Workshop on Recent Advances in Intrusion Detection, Cham, 2014
- [26] Ming, J., Xu, D., Jiang, Y., *et al.*: 'Binsim: trace-based semantic binary diffing via system call sliced segment equivalence checking'. Proc. of the 26th USENIX Security Symp., Vancouver, Canada, 2017
- [27] Blazytko, T., Contag, M., Aschermann, C., *et al.*: 'Syntia: synthesizing the semantics of obfuscated code'. 26th USENIX Security Symp., Vancouver, Canada, 2017